

Project Report

Statistical Methods for Machine Learning: Cats vs Dogs Binary Classification

LUCA DONGHI
luca.donghi1@studenti.unimi.it, 982862

30/03/2023

1. Task

"Use Tensorflow 2 to train a neural network for the binary classification of cats and dogs based on images from *this dataset*. Images must be transformed from JPG to RGB (or grayscale) pixel values and scaled down. Experiment with different network architectures and training parameters documenting their influence on the final predictive performance. Use 5-fold cross-validation to compute your risk estimates. While the training loss can be chosen freely, the reported cross-validated estimates must be computed according to the zero-one loss."

2. Introduction

The binary classification of cat and dog images is a classic beginner project in machine learning that helps to understand the functioning of these models. This report will explain every part of the project, from the data preprocessing to the computation of the risk estimates of the final model. To accomplish the task, several Convolutional Neural Network (CNN) architectures have been trained, and each of them has been evaluated and compared to previous models to modify the models and their hyperparameters accordingly.

This research, and consequently, this report, is structured as follows. Section 3 describes the provided data, its organization, and preprocessing. Section 4 presents the model architectures and the detailed process used to arrive at the final model, including hyperparameter tuning and model selection. Section 5 evaluates the final model, showcasing the types of errors it makes and presenting the cross-validated risk estimate. Finally, Section 6 concludes the report by presenting essential observations and lessons learned throughout the project.

3. Data

The *CatsVSDogs* dataset comprises 25,000 images, including 12,500 cat images and 12,500 dog images. The images have varying sizes, qualities, and orientations.

3.1. Data organization

The entire dataset was utilized for training, validating, and testing the models. The original dataset comprised two folders containing cat and dog images, respectively. To preprocess the images for input into the CNN models, a new folder named "data" was created that contained all 25,000 images. To preserve the label information, the images were renamed before merging them into the new folder, with a '0' prefix added to the cat images and a '1' prefix to the dog images. While loading the data into the coding environment two images were discarded as they were corrupted and they would have created problems later on. Then, the data was split into training, validation, and test sets using the `train_test_split` function from scikit-learn, with a test size of 20% and a validation size of 25% of the remaining data. (Training set = 10498 images, Validation set = 5000 images, Test set = 5000 images).

3.2. Preprocessing

The preprocessing involved three fundamental operations. First, all the images were reshaped into the commonly used size of 150x150 to avoid computational and memory capacity issues. Secondly, the images were converted into NumPy arrays with RGB channels. Even though grayscale conversion would have speeded up the computation, it was avoided as I believe color plays an important role in classifying cats and dogs. Finally, pixel values were scaled down between 0 and 1. Normalizing images in the pre-processing stage for CNNs is crucial to improve training convergence, avoid numerical problems, and make the features more homogeneous, ultimately enhancing the learning and generalization of the model.

4. Algorithm implementation

In this section, I will describe the model architectures and the process used to select them. To achieve this, I will use Convolutional Neural Network (CNN), which is a type of neural network that utilizes convolutional layers to extract features from images. The convolutional layer convolves a filter over the pixels of the input image to detect patterns and extract features. With each subsequent convolutional layer, the network can detect more complex features. In addition to convolutional layers, pooling layers (usually max pooling) are added to reduce the dimensionality of the feature maps and decrease the number of parameters to learn. CNNs are highly effective for image recognition tasks as they can learn to recognize complex features in images. This makes

them a suitable choice for classifying cats and dogs. Therefore, I will use CNNs to modify the structure of the models and their hyperparameters. By selecting the appropriate architecture and hyperparameters, the models can extract relevant features from the images and classify them accurately. Therefore, the selection process will be focused on finding the optimal combination of CNN layers, filter sizes, activation functions, and regularization techniques.

4.1. Model Selection

During the model selection process, different architectures and hyperparameters were experimented with, and each adjustment was based on the previous model's performance. A total of 10 models were trained using two metrics: binary cross-entropy as the loss function and accuracy as a measure of how well the model can classify input images. The loss function measures the difference between the predicted probability and the actual class label. On the other hand, accuracy indicates the percentage of correctly classified images. By monitoring both metrics, we can ensure that the model is not only minimizing the loss but also achieving high accuracy, which is crucial for the model to generalize well to new and unseen data.

- **Model1**

The *model1* serves as the baseline for the analysis as it is a very simple convolutional neural network. Its architecture is structured as follows:

- **Conv2D layers:** These are the convolutional layers that learn to extract features from the input images. The first layer has 32 filters, while the following layers have 64 and 128 filters. The filter size is (3, 3) for all layers, and the activation function used is *Relu*, which introduces nonlinearity and helps the network learn more complex patterns.
- **MaxPooling2D layers:** These layers perform downsampling, reducing the spatial dimensions of the feature maps. This helps reduce the number of parameters in the model and thus reduces the computational cost, while also improving the model's ability to recognize features in different scales and translations.
- **Flatten layer:** This layer reshapes the output of the previous *MaxPooling2D* layer into a one-dimensional tensor, so it can be fed into the dense layers.
- **Dense layers:** These are fully connected layers that perform classification based on the extracted features. The first dense layer has 512 units with *Relu* activation. The last dense layer has only one unit with *Sigmoid* activation, which outputs the probability of the image being a dog (dog=1).
- **Model compilation:** The model is compiled using the *Adam* optimizer, which is an adaptive learning rate optimization algorithm. It adjusts the learning rate during training, making it a popular choice for training neural networks.




Figure 1. Model1 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

As we can see from Image 1, *model1* exhibits significant overfitting as it achieves a high training accuracy of 97.9% but a much lower validation accuracy of 83.74%. The increasing validation loss after the 6th epoch, while training loss continues to decrease, indicates the model is fitting the training data too closely and struggles to generalize to new, unseen data. Addressing overfitting is crucial to improve the model's performance on real-world data.

- **Model2**

The *model2* adds a Dropout layer, which combats overfitting by randomly disabling 50% of input units during training, encouraging the model to learn more generalizable features.




Figure 2. Model2 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

As we can see in Image 2 *model2* has a lower training accuracy (71.44%) and validation accuracy (73.34%) compared to *model1*. The smaller gap between the training and validation accuracies indicates a reduced overfitting issue. The *model2* is likely to be underfitting the data and I need to find an optimal balance in this tradeoff.

• Model3

The *model3* is similar to the second model but modifies the Dropout layer. Instead of dropping 50% of input units during training, it drops 30%. This change aims to strike a balance between preventing overfitting and retaining enough information for the model to learn effectively.




Figure 3. Model3 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The *model3* demonstrates better performance compared to the previous models. The validation accuracy reached 87.56% while the training accuracy reached 96.27%, indicating a smaller gap between the two than in *model1*. However, it's important to note that starting from Epoch 11, the validation loss starts to increase, indicating that the model is still overfitting.

• Model4

model4 introduces several changes compared to the previous model:

- **Additional Dropout layers:** In *model4*, Dropout layers are added after each *Conv2D*–*MaxPooling2D* block, with dropout rates of 10%, 20%, and 30% respectively. These additional Dropout layers help to further reduce overfitting.
- **Early Stopping:** An *EarlyStopping* callback is used during training, monitoring the validation

loss with a *patience* of 3 epochs. This means that the training will stop if there is no improvement in validation loss for 3 consecutive epochs. The best weights, corresponding to the lowest validation loss, will be restored to the model at the end of training.




Figure 4. Model4 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The model is struggling to learn the underlying patterns in the data as evidenced by the slow decrease in training loss and high and inconsistent validation loss. Although there is no overfitting, the poor overall performance suggests that the model is underfitting the data: worst performance until now.

• Model5

The *model5* architecture introduces several changes compared to *model4*:

- **Batch Normalization:** In this new model, Batch Normalization layers are added after each *Conv2D* layer and before the Dense layer. These layers help normalize the inputs to each layer, which can speed up the training process and improve generalization.
- **Activation functions:** Instead of including the activation functions directly within the *Conv2D* and *Dense* layers, they are now separated into individual Activation layers. This allows Batch Normalization to be applied before the activation functions.
- **Dropout rate:** The dropout rate has been changed to a constant value of 0.2 for all Dropout layers.
- **Learning rate:** The learning rate for the *Adam* optimizer has been set to 0.0005, which is a smaller value compared to the default learning rate. This can help the model converge more steadily, potentially leading to better performance.




Figure 5. Model5 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

As we can see looking at Image 5 *model5* is back on track! However, even if it has shown a huge improvement in performance compared to *model4* it is not the best model I trained so far. Its training loss decreases consistently over epochs, while the validation loss increases after epoch 7. This model exhibits overfitting as the training accuracy is much higher than the validation accuracy and there is space for improvements.

• Model6

The *model6* again introduces several modifications compared to the previous model:

- **Reduced dropout rate:** The dropout rate is reduced to 0.1.
- **L2 regularization:** L2 regularization (with a lambda of 0.001) is added to the *convolutional* and *dense* layers, which helps prevent overfitting by penalizing large weights in the model.
- **Data augmentation:** The training data is augmented using rotation, width and height shifts, horizontal flips, and zoom. This increases the diversity of the training data and helps the model generalize better to unseen data.
- **Increased dropout rate after third convolutional layer:** A higher dropout rate of 0.5 is applied after the third convolutional layer to reduce overfitting on deeper layers.




Figure 6. Model6 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The *model6* shows consistent improvement in both training and validation accuracy compared to previous models. Any overfitting is minimal, thanks to the use of regularization techniques and data augmentation. It's also worth noting that validation accuracy is sometimes higher than training accuracy due to regularization techniques applied only during training and data augmentation.

• Model7

The *model7* differs from the previous model by adding two more *convolutional* layers, making the architecture deeper. It also includes an additional *Dense* layer with 1024 units. These changes aim to improve the model's capacity to learn more complex and hierarchical features.




Figure 7. Model7 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The *model7* showed a significant improvement in validation accuracy compared to the previous model, reaching 88.18% at epoch 12. This improvement can be attributed to the deeper architecture and the additional Dense layer with 1024 units, which allowed the model to learn more complex and hierarchical features. However, there was a slight increase in training time due to the increased model complexity. Despite the increase in validation accuracy, we still observe some fluctuation in the validation accuracy and loss during training.

• Model8

The *model8* introduces a *GlobalAveragePooling2D* layer instead of the *Flatten* layer used in the previous model. Global average pooling reduces the spatial dimensions of the feature maps by computing the average value of each channel, resulting in a smaller and more compact representation. This change helps reduce the total number of parameters, which can help prevent overfitting and improve training efficiency.




Figure 8. Model8 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The *model8* displays a consistent improvement in validation accuracy, peaking at 89.10% in epoch 13. When compared to *Model7*, which achieved a peak validation accuracy of 88.18% at epoch 12, *Model8* demonstrates better performance. However, like *model7*, *model8* also experiences fluctuations in validation accuracy and loss, indicating potential benefits from further fine-tuning or additional regularization techniques for more stable performance.

• Model9

The *model9* introduces several changes compared to the previous one:

- **Reduced dropout rate:** The *dropout_rate* is reduced to 0.1
- **Lower learning rate:** The *learning_rate* is lowered to 0.0005, which may help the model converge more smoothly and find better local minima.
- **L2 regularization:** It is added to the *convolutional* and *dense* layers with an *l2_lambda* value of 0.001.
- **Increased batch size:** The *batch_size* is increased to 64, which can provide a better estimate of the gradient and potentially accelerate training.




Figure 9. Model9 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The *model9* shows a clear improvement over *model8*. The highest validation accuracy achieved by *Model9* is 93.20% at epoch 23, compared to *Model8*'s peak validation accuracy of 89.10% at epoch 13. This indicates that the adjustments made in *Model9* have effectively improved the model's performance in capturing the underlying patterns in the data. Moreover, *Model9* appears to have a relatively more stable performance, with smaller fluctuations in validation accuracy and loss, thanks to the inclusion of L2 regularization and the reduced dropout rate. This suggests that *Model9* has a better balance between learning capacity and generalization, leading to a more robust and accurate model. However, as the training didn't stop because of the early stopping callback, I think that the model could do even better if trained for more epochs.

• Model10

In *model10*, the number of epochs has been increased from 25 to 60, allowing the model more time to learn from the dataset. This change aims to explore whether additional training time could further improve the *model9*'s performance. All other aspects of *model10* remain the same as in *model9*, including the architecture, L2 regularization, dropout rates, and data augmentation.




Figure 10. Model10 Training and Validation Metrics. The left figure shows the accuracy while the right one shows the loss.

The final model, *model10*, demonstrates strong overall performance in terms of accuracy and loss. It achieved a validation accuracy of 94.08% at epoch 29, which shows that longer training improved its results, as expected. Compared to previous models, *model10*'s validation accuracy consistently increased or remained relatively stable over the epochs. Its accuracy improved gradually over the epochs, without any sign of overfitting. Based on these characteristics, *model10* is a satisfying final model for this task.

5. Final model evaluation

In this Section, I will evaluate *model10* performances on the test set, analyze the misclassified images, and present the cross-validated risk estimates.

The test set metrics for *model10* are:

- **loss:** 0.2293

- **accuracy:** 0.9418

94% of the test set images have been correctly classified.

5.1. Misclassified images

In Figure 11, the confusion matrix show that the performance is pretty balanced: the proportion of mistakes on dog images and on cat images is almost the same.




Figure 11. Model10 Confusion Matrix.

In Figure 12, a random sample of images that were misclassified by *model10* on the test set is shown. While some mistakes may be due to difficult images with poor quality or unusual positioning, it's also apparent that there are some relatively straightforward images for which the classification probabilities of being a cat or a dog are very close. This suggests that there is potential for improvement and that with more time and computational resources, it may be possible to train a better model.

5.2. Model10 cross-validation

Finally, I used 5-fold cross-validation to compute the risk estimates of the final model. As required by the task, even if I used the binary cross-entropy to train all the models, I used the zero-one loss for the cross-validation.

Zero-One Loss	<i>Model10</i>
Fold 1	0.0555
Fold 2	0.0407
Fold 3	0.0420
Fold 4	0.0383
Fold 5	0.0390
Mean	0.0431

Table 1: Model10 Cross-Validated Risk Estimate.

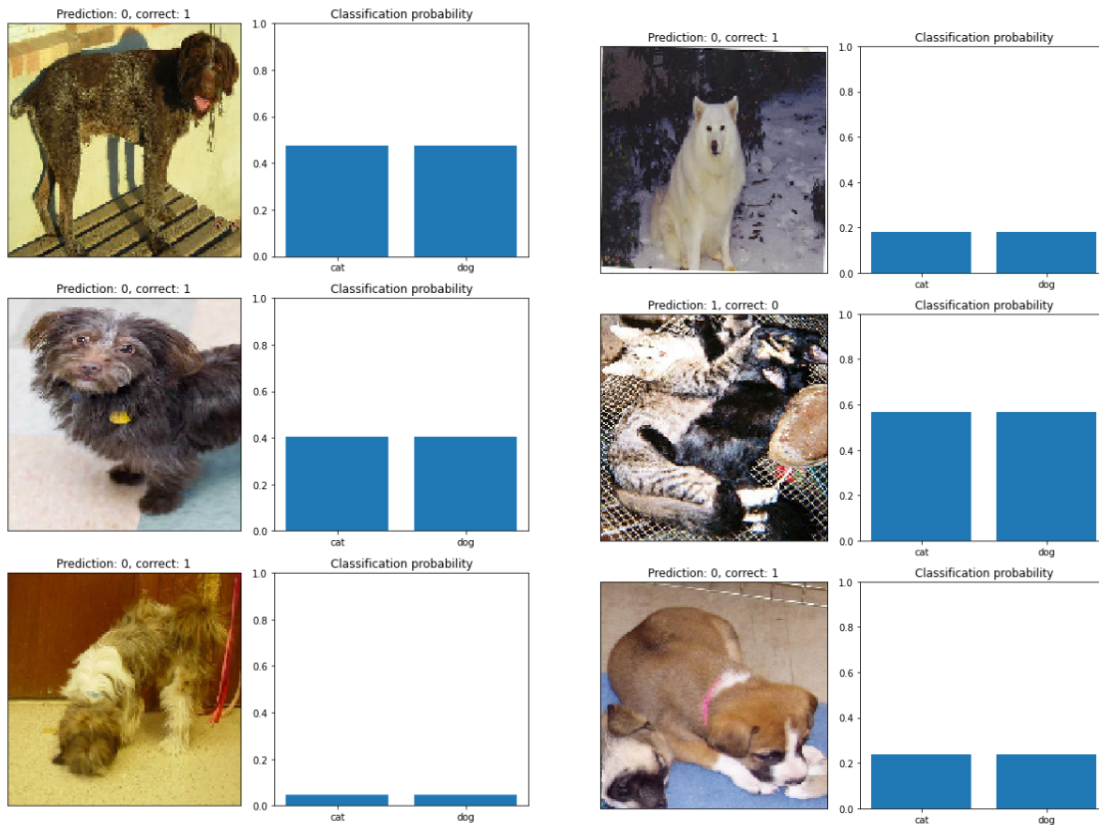


Figure 12. Model10 test set misclassified images. Each misclassified image is associated with the estimated probability of being classified as a dog or as a cat.

As shown in Table 1 the final cross-validated risk estimate of *model10* is 0.0431.

6. Conclusions

In this report, we have explored the development of a deep learning model for the classification of cats and dogs in images. We started by training a basic model and gradually introducing improvements to its architecture, regularization techniques, and data augmentation. After evaluating several models, we arrived at a final model, *model10*, which achieved satisfying 94% accuracy on the test set and 0.0431 as cross-validated risk estimate.

While the final model showed strong overall performance, we also analyzed misclassified images and cross-validated risk estimates to identify potential areas for improvement. The analysis showed that while the model performed well overall, there were still some challenging images where it struggled to correctly classify the subject. These results suggest that with more time and computational resources, it may be possible to train even better models that perform more consistently across a wider range of images.

I declare that this material, which I now submit for assessment, is entirely my own work and has not been taken from the work of others, save and to the extent that such work has been cited and acknowledged within the text of my work. I understand that plagiarism, collusion, and copying are grave and serious offences in the university and accept the penalties that would be imposed should I engage in plagiarism, collusion or copying. This assignment, or any part of it, has not been previously submitted by me or any other person for assessment on this or any other course of study.